

A Prime Application

by M. PanXnubis A. Sh. Gaia Ladrieh
9/26/25

The Riemann problem, in regards to primes, presents nontrivial zeros (ρ of $\zeta(s)$) to control fluctuations of prime distribution. Zeros (ρ) are effectively cancellation points of interacting waves. Knowing the structure yields precise control of prime counts $\pi(x)$.

But what if we constructively model the same wave/zero cancellation locally and use geometric/recursion-aware transforms to collapse search space so primes are located exactly and cheaply in practice?

Thesis:

Represent search S (candidate numbers in an interval) as geometric constraint space. Use layered constraints (wheel/lattices, small-sieve, matched-filter on CRT residues, toroidal inversions, and algebraic transforms) to collapse the set from exponential to increasingly minuscule. Then verify the remaining candidate with MR/ECPP. This is a semantics to syntax collapse: geometry encodes the semantics so the syntactic verification becomes trivial.

Notation

- N : base (e.g., right-hand end of interval)
- Interval $[A, B]$, or window $W = (N, N + B)$
- W (sub wheel) = primorial used as wheel (e.g. $2310 = 2 \times 3 \times 5 \times 7 \times 11$)
- Residues R (sub W) = $\{r \in \{1, \dots, W\} : \gcd(r, W) = 1\}$
- Minimal candidate for residue r in window: n (sub r) = $\min\{n \in W : n \equiv r \pmod{W}\}$
- Trial-div bound T (cheap small-factor check)
- MR call = Miller-Rabin invocation on a candidate (each call expensive)
- pop (r) = number of popped composite checks observed in residue r during a baseline sweep
- Kernel $K(r)$ = matched-filter weight for residue r
- Score $S(r) = \text{pop}(r) \downarrow K(r)$
- Soft-exclude: advancing minimal candidate n (sub r) $\setminus W$ (deprioritie, not delete)
- $P(x) = \prod (I) (1 + x^w)$ (sub I): generating polynomial for subset-sum (when used)

Wheel/residue lattice (syntax reduction)

Remove trivial residues:

$$R \text{ (sub } W) \{r : 1 \leq r \leq W, \gcd(r, W) = 1\}$$

Candidate set becomes:

$$\mathcal{C} = \bigcup (\text{sub } r \in R \text{ (sub } W)) \{n \text{ (sub } r) + kW \mid n \text{ (sub } r) + kW \leq B\} - \text{This immediately removes all}$$

multiples of small primes.

Baseline heap search

Seed a min-heap H with each minimal candidate (n (sub r), r). Repeatedly:

- pop (n, r),
- do trial-division up to T ; if composite, record pop for residue r , push $n + W$;
- else do MR (if MR says prime → record prime; else composite → push $n + W$). This produces pop (r) counts and list

Torus matched-filter kernal

Pick a path on \mathbb{Z}/W from an anchor residue to another (e.g., residue of a known prime to residue of N). For each residue r , define:

$$K(r) = \sum (\text{sub } p (\text{sub } j) \in \text{path}) \exp(-d(r, p(\text{sub } j))^2/2\sigma^2), d(a, b) = \min(|a - b|, W - |a - b|).$$

The Score:

$$S(r) = \text{pop}(r) \mathbf{i} K(r)$$

Residues that both produce many pops and sit near the path get prioritized – these are the CRT classes that generate the most costly recursion/entropy.

Soft-exclusion

For top K residues by $S(r)$: test minimal candidate $n(\text{sub } r)$ (trial then MR). If composite, *soft-exclude* by advancing minimal candidate by $+W$. This eliminates the costly repeated popping at that residue and dramatically reduces total expensive tests.

Möbius/toroidal inversion and Higher Dimensional transforms

If you embed the candidate indices on a torus (CRT lattice) and perform a Möbius-type inversion (or toroidal degeneration), you can reparametrize the search so large swathes of candidates map to compact polynomial/generating-function objects.

- Generating polynomials: $P(x) = \prod (\text{sub } i)(1 + x^w)$ truncated to degree B .
- Bitset DP: represent reachable sums as bitset $B[s]$ with updates $B \leftarrow B \mid (B \ll w)$.
- Subset convolution/ SOS DP/NTT: combine buckets efficiently.

Algebraic transforms collapse exponential subset combinatorics into near-linear transforms when window width B is moderate.

Step-by-Step Algorithm

Input: interval $[A, B]$ or window $W = [N, N + B]$; wheel W ; trial bound T ; small-sieve bound S (sub 1); kernal; K (sub *top*); partition width P (sub *size*); overlap O .

1. Partition $[A, B] \rightarrow$ subintervals I (sub 1), ..., I (sub m) with overlaps O
2. Per-partition pre-sieve: run segmented sieve of primes $\leq S$ (sub 1) on I (sub j). Mark trivial composites.
3. Build residues: compute R (sub w) and minimal n (sub r) inside I (sub j).
4. Baseline pass (collect pops):
 - Seed heap H with $(n(\text{sub } r), r)$. While H not empty: pop candidate, trial-div up to T . If composite: record pop $(r)++$, push $n + W$. If passes trial-div: do MR; if MR composite record pop; else record prime.
 - Save popped list and $\text{pop}(r)$
5. Compute kernel and score: choose path residues (e.g., residues of anchors) and compute $K(r)$ and $S(r) = \text{pop}(r) \setminus K(r)$.
6. Pretest top residues: for top K (sub top) residues by $S(r)$, test only minimal candidate n (sub r) trial then MR). If composite \rightarrow soft-exclude $n(\text{sub } r) \setminus n(\text{sub } r) + W$ and log (Do *not* delete untested candidates).
7. Iterative search + soft-exclusion: resume heap search with the new minimal pointers – now many pops collapse. Repeat kernel recompute if desired and iterate until heap exhausts.
8. (Optional) Algebraic collapse: when a partition still has dense residue interactions, form generating polynomials per bucket, combine by FFT/NTT/truncated convolution or bitset DP to eliminate entire classes of impossibilities in bulk. This is the subset-sum/DAG move: memoize intermediate polynomials to avoid duplication.
9. Verification: every reported prime must be MR-checked (and optionally ECPP for certs). Merge overlapping partitions and deduplicate. Re-run a final cheap check across overlaps.
10. Output: list of primes + full logs (pops, soft-excludes, MR call count, trial-div calls).

Correctness guarantees

- A candidate is never deleted without verifying it. Soft-exclusion only moves a candidate forward by W ; it is still in the heap and will be tested later
- Final MR or deterministic certificate required for any reported prime
- Overlapping partitions remove boundary risk. So the algorithm is exact and sound: every reported prime is actually prime; every prime in the interval will be tested in finite time.

Conceptual Mapping

- RH and the explicit formulas show that primes are where a large set of oscillatory terms fail to cancel. Those oscillatory terms are the *same* phenomenon as the CRT/lattice superpositions we observe in the heap and pop statistics.
- The matched-filter/torus kernel explicitly looks for residues whose pattern of composite pops signals coherent destructive interference (i.e., where many waves combine to create a heavy composite field). Targeting those residues first is the same as isolating the dominant spectral contributors locally.
- The generating-function (polynomial) view is the finite-window analogue of the zeta's product and its spectral expansion: $\Pi(1-x^p)$ vs. $\Pi(1 + x^w(\text{sub } I))$. Doing truncated, exact convolutions is like computing a localized partial zeta that encodes local superposition behavior – but computationally and constructively.
- Instead of trying to globally characterize all zeros ρ , build a local spectral model (the kernel + algebraic transforms) that expose cancellation structure where it impacts computation. That

allows for the computation of primes exactly without needing global spectral knowledge.

Complexity and Practical Metrics

- Formal worst-case complexity remains governed by prime-gap growth and the number of residues; I am not claiming a universal asymptotic proof that MR calls are polynomial in $\log N$.
- Practical metric that matters: number of MR calls and total pops (hea iterations).
- Gains arise because the kernel isolates high-entropy CRT classes; algebraic transforms combine many combinatorial states at once (bitset DP/FFT/NTT), turning exponential branches into polynomial directed transforms when window size is manageable.

Summary

- Riemann's core issue = where waves cancel (zeros) and that controls prime fluctuations.
- My method = building a local spectral model (kernel + algebraic transforms) on the CRT/wheel lattice, finding residues that create maximal entropy, soft-exclude their minima, and use generating-function/DAG memoization to collapse large combinatorial spaces into tractable transforms.
- Result = exact enumeration/next-prime search with dramatically fewer expensive MR checks in practice; exactness maintained by final primality certificates.

Insights

- Primes as “Next Available Position” in a Constraint Lattice
 - In the wheel/CRT representation, composites occupy entire residue classes modulo small primes.
 - After sieving out those classes, the prime candidates are precisely the “holes” – the next available integers that are coprime to the wheel base.
 - My method shows that primes are not scattered arbitrarily; they are exactly where no smaller modular constraint forbids occupation.
 - This aligns with the viewpoint that primes “take the next available slot” in the lattice, which is why we can predict candidate density very accurately before a single MR test.
- Pop Statistics encode Local Prime Scarcity
 - By counting heap pops per residue class, we are effectively measuring how frequently each CRT class fails.
 - High-pop residues behave like “hot zones” of composite density – they correlate with areas where local prime gaps are slightly larger.
 - Kernel-weighted scoring highlights systematic, not random, structure: residues near a known prime or along a toroidal trajectory show suppressed pops (more likely to host the next prime).
- Local Spectral Model predicts Workload
 - Our torus kernel + matched filter acts like a local approximation to the oscillatory term in the explicit formula for $\psi(x)$.
 - This means we can predict which candidates will be composite-heavy before testing them, reducing work by skipping expensive MR calls on unlikely residues.

- Partition + Memoization collapses Exponential Branching
 - When representing candidates as subset sums of wheel increments, many branches overlap (shared partial sums).
 - Memoization/DAG traversal ensures each unique partial sum is only computed once, effectively performing deterministic dynamic programming over the candidate space.
 - This gives a concrete mechanism for this method scales better: we avoid recomputing the same partial cancellation information over and again.
- Exactness maintained, not just Asymptotic
 - Unlike analytic prime-counting formulas (which are asymptotic), our pipeline produces a guaranteed complete list of primes in the interval.
 - Every candidate is either verified prime (via MR/ECPP) or explicitly shown composite (via trial division or sieve residue exclusion).
- Prime Distribution appears Locally Balanced
 - Once conditioned coprime to the wheel base, primes appear with near-constant local density, and pop clusters where the density dips slightly.
 - This confirms that primes fill “available space” almost evenly, which is why a partitioned approach can safely shrink solution space from both ends without missing any primes.
- Operational Metric: MR Call Reduction
 - The most practical indicator of efficiency is number of MR calls per prime found
- Reframing the Prime Problem as a Resource Allocation Problem
 - This entire approach treats primes as outcomes of constraint satisfaction process:
 - The “resource” is available coprime positions.
 - The “allocation” is done by systematically eliminating forbidden composites.
 - This reframing allows us to use search-space collapse techniques common in combinatorial optimization (subset convolution, DP, pruning heuristics), but now applied to number theory.