Admissibility Gates for Constraint-Based Systems

A Pre-Verification Requirement for Meaningful Validation

by PanXnubis Gaia Ladrieh
1/3/26

1. Purpose

Verification and Validation (V&V) methodologies are widely employed to assess whether a system meets specified requirements and performs as intended under defined conditions. These methodologies presuppose, often implicitly, that the underlying system description is internally coherent, structurally identifiable, and jointly satisfiable.

This document asserts a prior requirement: **before verification or validation activities are meaningful, the system under evaluation must satisfy formal admissibility conditions**. Absent admissibility, verification results may be internally consistent yet externally meaningless, and validation outcomes may reflect artifacts of model construction rather than properties of the system itself.

The purpose of this document is to define a minimal, non-heuristic admissibility gate that determines whether a system's stated constraints can, in principle, be simultaneously true and evaluable. This gate operates *upstream* of verification and validation and does not replace them.

2. Scope

This document applies to systems characterized by explicit or implicit constraints, including but not limited to:

- engineered physical systems
- safety-critical and fault-tolerant architectures
- probabilistic and uncertainty-quantified models
- simulation-based and surrogate-assisted analyses
- data-driven and hybrid AI systems

The admissibility gate defined herein evaluates **structural properties of the system description**, not performance metrics, implementation fidelity, or empirical adequacy.

This document does **not**:

- assess correctness, optimality, or safety margins
- predict system behavior
- replace verification, validation, or certification processes
- introduce probabilistic assumptions or tolerances

Failure to satisfy admissibility does not imply system failure. It implies that **verification results cannot be interpreted as evidence of correctness.**

3. Definitions (Normative)

The following definitions are normative and apply throughout this document.

3.1 Constraint

A **constraint** is a declarative statement asserting a relationship, bound, requirement, or invariant that the system claims to satisfy. Constraints may be mathematical, logical, algorithmic, physical, or procedural.

Constraints are treated as *simultaneously asserted*, unless explicitly stated otherwise.

3.2 Admissibility

A system is **admissible** if and only if all explicitly stated constraints can be simultaneously satisfied under a single, non-contradictory interpretation.

Admissibility is binary. A system either satisfies admissibility or it does not.

3.3 Contradiction

A **contradiction** exists when two or more explicit constraints cannot be jointly true under the same system state or parameterization.

Resolution of contradiction through interpretation, prioritization, or informal reasoning is not permitted.

3.4 Identifiability

A system is **identifiable** if its stated constraints uniquely determine the quantities they purport to constrain, up to explicitly declared equivalence classes.

Undeclared decrees of freedom constitute non-identifiability.

3.5 Verification

Verification is the process of determining whether a system implementation conforms to its stated specifications and requirements.

Verification presumes that the specifications themselves are internally coherent and jointly satisfiable. Verification does not evaluate the admissibility of the specification set.

3.6 Validation

Validation is the process of determining whether a system adequately represents or performs relative to an intended real-world use or reference.

Validation presumes that the verified system corresponds to a meaningful, well-defined target. Validation does not resolve contradictions or non-identifiability present in the system description.

3.7 Observer Independence

An admissibility determination is **observer independent** if any competent evaluator, given the same explicit constraints and rules, arrives at the same admissibility outcome.

Admissibility evaluations shall not depend on:

- undocumented assumptions
- expert intuition
- contextual interpretation
- corrective reinterpretation of constraints

4. Admissibility Requirements

A system description SHALL be considered admissible **if and only if** all of the following requirements are satisfied.

These requirements are evaluated prior to verification or validation activities.

4.1 Explicit Constraint Enumeration

All constraints relevant to system behavior SHALL be explicitly enumerated.

Implicit assumptions, informal expectations, or context-dependent interpretations SHALL NOT be used to supplement the constraint set

during admissibility evaluation.

4.2 Simultaneity Requirement

All explicitly stated constraints SHALL be treated as simultaneously asserted.

Sequential activation, conditional relevance, or selective applicability of constraints MUST be explicitly declared. Absent such declaration, simultaneity is assumed.

4.3 Non-Contradiction Requirement

The constraint set SHALL admit at least one system state in which all constraints are jointly satisfied.

If no such state exists, the system is inadmissible.

Resolution of contradiction by:

- informal prioritization
- post hoc reinterpretation
- probabilistic dilution
- asymptotic reasoning

is not permitted.

4.4 Syntax Precedence Rule

Syntactic statements SHALL take precedence over semantic interpretation.

The meaning of a constraint SHALL be determined by its explicit form and declared variables. Interpretive reconciliation that alters syntactic content is prohibited.

If semantic interpretation conflicts with syntactic expression, the constraint set is inadmissible.

4.5 Identifiability Requirement

For each constrained quantity, the constraint set SHALL uniquely determine its admissible values, up to explicitly declared equivalence classes.

Undeclared free parameters, latent variables, or observer-selected degrees of freedom constitute non-identifiability and render the system inadmissible.

4.6 Independence from Verification Outcomes

Admissibility SHALL NOT be inferred from successful verification or validation results.

A system may pass verification and validation while remaining inadmissible.

4.7 Deterministic Outcome Requirement

Given the same explicit constraint set and admissibility rules, independent evaluators SHALL arrive at the same admissibility determination.

If admissibility depends on evaluator discretion, the system description is inadmissible.

5. Implications for Verification and Validation

An inadmissible system description undermines the interpretability of verification and validations results.

Specifically:

- Verification may confirm internal consistency relative to an incoherent specification.
- Validation may succeed by fitting artifacts rather than system properties.
- Safety margins may aggregate incompatible assumptions without detection.
- Uncertainty quantification may obscure structural contradiction through stochastic smoothing.

These outcomes do not indicate methodological failure; they indicate **misplaced application.**

6. Discussion

The admissibility gate defined herein does not impose additional performance requirements, accuracy thresholds, or domain-specific criteria. It imposes only logical and structural conditions that are already tacitly assumed by downstream processes.

The cost of admissibility evaluation is minimal. The cost of proceeding without it is unbounded.

7. Illustrative Failure Modes (Normative, Non-Exhaustive)

The following sections illustrate common classes of system

descriptions that fail admissibility requirements defined in Section 4. These examples are representative, not exhaustive, and are presented without reference to specific implementations or organizations.

7.1 Monte Carlo-Based Uncertainty Quantification

7.1.1 Description

Monte Carlo-based uncertainty quantification (UQ) is widely used to propagate uncertainty through complex systems by repeated sampling of input distributions and observing resulting output distributions.

Typical formulations involve:

 – probabilistic input parameterization,
 – stochastic sampling over assumed distributions,
 – aggregation of outputs into confidence intervals, exceedance probabilities, or risk metrics

These methods are commonly applied in safety, performance, and reliability analysis.

7.1.2 Constraint Structure

In such systems, the explicit constraint set typically includes:

 – bounds on input parameters,
 – assumed probability distributions,
 – model equations mapping inputs to outputs,
 – acceptance criteria defined in terms of statistical measures.

However, the following elements are frequently implicit or under-specified:

 – the admissibility of assumed distributions,
 – the simultaneity of constraints across samples,
 – the identifiability of latent or unobservable parameters,
 – the consistency of constraints under joint realization.

7.1.3 Admissibility Evaluation

Under Section 4, admissibility requires that all explicit constraints be jointly satisfiable and that constrained quantities be identifiable.

In Monte Carlo UQ formations:

- constraints are evaluated *sample-wise*, not jointly,
  - contradiction between constraints may be distributed across samples rather than resolved,
  - non-identifiability is masked by distributional assumptions,
  - inadmissible regions are integrated rather than rejected.

Probabilistic aggregation does not eliminate contradiction; it obscures it.

Resolution of inconsistency by appeal to asymptotic convergence, large sample size, or distributional smoothing constitutes semantic reinterpretation and is prohibited under Section 4.4

7.1.4 Outcome

Because the constraint set does not admit a jointly satisfiable realization independent of sampling procedure, and because identifiability is assumed rather than demonstrated, the system description fails admissibility requirements.

Successful statistical convergence does not establish admissibility.

7.1.5 Determination

Therefore, Monte Carlo-based uncertainty quantification systems, as commonly specified, are inadmissible under Section 4.

7.2 Safety Margin Aggregation (Load and Resistance Factor-Style Formulations)

7.2.1 Description

Safety margin aggregation methods are commonly employed in engineering and safety-critical domains to ensure acceptable performance under uncertainty. Representative approaches include load and resistance factor design (LRFD) and related margin-based formulations.

These approaches typically:

  - decompose system behavior into contributing loads and resistances,
  - apply partial safety factors to individual terms,
  - aggregate factored quantifies into a global safety margin or utilization ratio,
  - declare acceptability when the aggregated margin exceeds a prescribed threshold

7.2.2 Constraint Structure

The explicit constraint set in such formulations usually includes:

- nominal load models,
- nominal resistance models,
- safety factors applied to each component,
- an acceptance inequality defining adequate margin.

However, these formulations commonly leave implicit:

- the joint realizability of factored extremes,
- the simultaneity of worst-case assumptions,
- the dependency structure among contributing terms,
- the identifiability of the governing failure state.

Safety factors are often calibrated independently and combined additively or multiplicatively without explicit declaration of joint admissibility.

7.2.3 Admissibility Evaluation

Under Section 4, all explicit constraints must be jointly satisfiable and evaluated simultaneously.

In safety margin aggregation:

- factored loads and factored resistances are treated as if jointly realizable without demonstration,
- incompatibilities between worst-case assumptions are absorbed into margin arithmetic,
- constraint violation is redistributed rather than resolved,
- admissibility is inferred from margin size rather than constraint coherence.

The aggregation of independently conservative assumptions does not establish the existence of a realizable system state satisfying all constraints.

Appeal to historical calibration, code compliance, or empirical conservatism constitutes semantic justification and does not satisfy the syntactic requirements of admissibility.

7.2.4 Outcome

Because the aggregated safety margin does not correspond to a demonstrably admissible joint system state, and because the governing

constraint realization is not identifiable, the specification fails
admissibility requirements.

Passing a margin check does not imply the existence of a coherent
underlying system configuration.

7.2.5 Determination

**Therefore, safety margin aggregation-based system descriptions, as
commonly specified, are inadmissible under Section 4.**

7.3 Redundant and Fault-Tolerant Systems (Including Diverse Redundancy)

7.3.1 Description

Redundant and fault-tolerant system architectures are widely employed
to improve reliability, availability, and safety. Such systems
incorporate multiple components, subsystems, or pathways intended to
compensate for individual failures.

Representative strategies include:

 - parallel redundancy,
 - voting schemes,
 - failover mechanisms,
 - diverse redundancy using heterogeneous implementations to
   mitigate common-mode failure.

These architectures are prevalent in aerospace, nuclear,
infrastructure, and safety-critical software systems.

7.3.2 Constraint Structure

The explicit constraint set in fault-tolerant system descriptions
typically include:

 - functional requirements for each redundant element,
 - failure assumptions for individual components,
 - fault-detection and isolation logic,
 - system-level performance or safety requirements under fault
   conditions.

Frequently implicit or under-specified are:

 - the joint failure space across redundant elements,
 - the simultaneity of failure and recovery assumptions,
 - the consistency of behavioral assumptions across diverse

implementations,
    – the identifiability of the system's governing failure mode.

Redundancy is often treated as a structural property rather than as a
constraint set requiring joint satisfaction.

7.3.3 Admissibility Evaluation

Under Section 4, admissibility requires that all explicit constraints
be simultaneously satisfiable and that constrained behaviors be
identifiable.

In redundant and fault-tolerant systems:

    – component-level correctness is evaluated independently rather
      than jointly,
    – failure scenarios are partitioned rather than composed,
    – incompatibilities between fault assumptions are masked by
      architectural layering,
    – diversity is assumed to imply independence without explicit
      constraint declaration

The existence of multiple pathways does not demonstrate the existence
of a coherent global state system satisfying all stated constraints.

Appeal to architectural robustness, defense-in-depth, or qualitative
diversity arguments constitutes semantic interpretation and is
prohibited under Section 4.4.

7.3.4 Outcome

Because the constraint set does not admit a uniquely identifiable,
jointly satisfiable system state across failure conditions, and
because contradiction is absorbed rather than rejected, the system
description fails admissibility requirements.

Successful fault tolerance under selected scenario does not establish
admissibility of the overall specification.

7.3.5 Determination

**Therefore, redundant and fault-tolerant system descriptions,
including those employing diverse redundancy, are inadmissible under
Section 4.**

7.4 Data-Driven and AI-Based Model Validation

7.4.1 Description

Data-driven and AI-based models are increasingly employed in prediction, classification, control, and decision-support roles. Validation of such models typically relies on empirical performance measures derived from data, including:

- training and test set accuracy,
- cross-validation error,
- out-of-distribution performance metrics,
- robustness or stress-testing under perturbed inputs.

These approaches are frequently presented as substitutes for, or supplements to, physics-based or mechanistic modeling.

7.4.2 Constraint Structure

The explicit set in AI-based system descriptions generally includes:

- a model architecture or hypothesis class,
- an optimization objective,
- a training dataset and evaluation protocol,
- acceptance criteria defined in terms of empirical performance.

Commonly implicit or under-specified are:

- the admissibility of the data as a representative constraint carrier,
- the simultaneity of learned behaviors across operating regimes,
- the identifiability of causal or governing variables,
- the stability of constraints under distributional shift.

Performance metrics are often treated as proxies for system validity rather than as constraints requiring joint satisfaction.

7.4.3 Admissibility Evaluation

Under section 4, admissibility requires that all explicit constraints be simultaneously satisfiable and that constrained quantities be identifiable.

In data-driven validation:

- constraints are induced statistically rather than declared syntactically,

- contradictions between learned behaviors are averaged rather than resolved,
- non-identifiability is absorbed into latent representations,
- success is defined by empirical fit rather than by constraint coherence.

Improved predictive performance does not demonstrate that a coherent system description exists. It demonstrates only that a function approximation has been tuned to observed data.

Appeal to generalization, scale, or benchmark dominance constitutes semantic justification and does not satisfy admissibility requirements.

7.4.4 Outcome

Because the system description does not uniquely identify governing variables or admit a jointly satisfiable constraint realization independent of data selection, the specification fails admissibility requirements.

Passing validation benchmarks does not establish admissibility of the system description.

7.4.5 Determination

Therefore, data-driven and AI-based model validation approaches, as commonly specified, are inadmissible under Section 4.

8. Procedural Integration in Verification and Validation Workflows

8.1 Placement of the Admissibility Gate

The admissibility evaluation defined in this document SHALL be performed **prior** to verification and validation activities.

The admissibility gate is upstream of:

- requirements verification,
- model verification,
- validation against data or reference systems,
- uncertainty quantification,
- safety margin evaluation.

Failure at the admissibility stage precludes meaningful interpretation of downstream results.

8.2 Relationship to Existing Practice

The admissibility gate does not replace verification or validation processes. It constrains their domain of applicability.

Specifically:

- Verification SHALL be applied only to admissible specifications.
- Validation SHALL be interpreted only for admissible systems.
- Statistical or architectural success SHALL NOT be used to infer admissibility retroactively.

Existing V&V methodologies remain valid within their proper scope.

8.3 Implementation Characteristics

An admissibility evaluation:

- requires no simulation,
- requires no empirical data,
- requires no performance thresholds,
- requires no domain-specific calibration.

It requires only:

- explicit constraint declaration,
- simultaneous evaluation,
- syntactic consistency,
- identifiability determination.

The evaluation is deterministic and observer independent.

8.4 Consequences of Non-Adoption

Proceeding without admissibility evaluation permits:

- verification of incoherent specifications,
- validation of non-identifiable systems,
- aggregation of incompatible assumptions,
- masking of contradiction through statistical or architectural means.

These outcomes are not detectable downstream.

8.5 Summary

Admissibility is a prerequisite, not a refinement.

Verification and validation presuppose what admissibility establishes.

Failure to evaluate admissibility constitutes a category error.

Appendix A: Formal Admissibility Gate (Normative)

A.1 Constraint Set

Let

$C = \{c\_1, c\_2, …, c\_n\}$

be the set of explicitly stated constraints defining a system description.

Each constraint c_i is a syntactic statement over declared variables.

A.2 Admissibility Condition

The system description is **admissible** *if and only if*:

$∃x∈X$ such that $⋀\_(i=1)(\text{upper } n)\ c\_i(x)$ = true

where:

- X is the declared state space,
- all constraints are evaluated simultaneously,
- no reinterpretation ore relaxation is permitted.

A.3 Contradiction Gate

If:

$∀x∈X,\ ∃c\_i∈C$ such that $c\_i(x)$ = false

then the system description is inadmissible.

A.4 Identifiability Requirement

For each constrained quantity q, the constraint set SHALL uniquely determine admissible values:

$|\{q \mid C(q)\}|$ = 1

up to explicitly declared equivalence classes.

Undeclared degrees of freedom imply non-identifiability and inadmissibility.

A.5 Syntax Precedence

Let $c_i(syn)$ denote syntactic form of constraint c_i and c_i(sem) any semantic interpretation.

If:

$c_i(sem) \neq c_i(syn)$

then $c_i(syn)$ SHALL govern.

Semantic reconciliation that alters syntax is prohibited.

A.6 Determinism

Given identical C and evaluation rules, all competent evaluators SHALL reach the same admissibility determination.

If not, the specification is inadmissible.

A.7 Note on Implementation

One admissibility evaluation kernel implementing this gate is described elsewhere (MidnightSun). The admissibility criteria defined herein are independent of any specific implementation.